

Subsetting Vectors

Raphael Carvalho

20/05/2019

Subsetting Vectors

The way you tell R that you want to select some particular elements (i.e. a ‘subset’) from a vector is by placing an ‘index vector’ in square brackets immediately following the name of the vector. For a simple example, try `x[1:10]` to view the first ten elements of `x`.

```
x[1:10]
```

```
## [1]          NA          NA  1.042252547          NA          NA
## [6] -0.092668232 -0.500333319 -1.404759367          NA  0.008206288
```

What do you think `x[is.na(x)]` will give you?

```
print("2: A vector of all NAs")
```

```
## [1] "2: A vector of all NAs"
```

Recall that `!` gives us the negation of a logical expression, so `!is.na(x)` can be read as ‘is not NA’. Therefore, if we want to create a vector called `y` that contains all of the non-NA values from `x`, we can use `y <- x[!is.na(x)]`. Give it a try.

```
y <- x[!is.na(x)]
```

Recall that the expression `y > 0` will give us a vector of logical values the same length as `y`, with `TRUE`s corresponding to values of `y` that are greater than zero and `FALSE`s corresponding to values of `y` that are less than or equal to zero. What do you think `y[y > 0]` will give you?

```
print("1: A vector of all the positive elements of y")
```

```
## [1] "1: A vector of all the positive elements of y"
```

You might wonder why we didn’t just start with `x[x > 0]` to isolate the positive elements of `x`. Try that now to see why.

```
x[x>0]
```

```
## [1]          NA          NA  1.042252547          NA          NA
## [6]          NA  0.008206288          NA          NA          NA
## [11]          NA          NA          NA  0.278735977          NA
## [16]          NA          NA          NA  1.610933847  0.369899514
## [21]          NA  0.043526480          NA          NA          NA
## [26]          NA  0.428470483  0.846031063
```

Since `NA` is not a value, but rather a placeholder for an unknown quantity, the expression `NA > 0` evaluates to `NA`. Hence we get a bunch of `NAs` mixed in with our positive numbers when we do this.

Combining our knowledge of logical operators with our new knowledge of subsetting, we could do this – `x[!is.na(x) & x > 0]`. Try it out.

```
x[!is.na(x) & x > 0]
```

```
## [1] 1.042252547 0.008206288 0.278735977 1.610933847 0.369899514 0.043526480
## [7] 0.428470483 0.846031063
```

Can you figure out how we'd subset the 3rd, 5th, and 7th elements of `x`? Hint – Use the `c()` function to specify the element numbers as a numeric vector.

```
x[c(3, 5, 7)]
```

```
## [1] 1.0422525      NA -0.5003333
```

It's important that when using integer vectors to subset our vector `x`, we stick with the set of indexes $\{1, 2, \dots, 40\}$ since `x` only has 40 elements. What happens if we ask for the zeroth element of `x` (i.e. `x[0]`)? Give it a try.

```
x[0]
```

```
## numeric(0)
```

As you might expect, we get nothing useful. Unfortunately, R doesn't prevent us from doing this. What if we ask for the 3000th element of `x`? Try it out.

```
x[3000]
```

```
## [1] NA
```

Luckily, R accepts negative integer indexes. Whereas `x[c(2, 10)]` gives us ONLY the 2nd and 10th elements of `x`, `x[c(-2, -10)]` gives us all elements of `x` EXCEPT for the 2nd and 10 elements. Try `x[c(-2, -10)]` now to see this.

```
x[c(-2, -10)]
```

```
## [1]      NA 1.0422525      NA      NA -0.09266823
## [6] -0.50033332 -1.40475937      NA -0.33069581 -1.24412832
## [11]      NA      NA      NA -0.40215828      NA
## [16]      NA -1.49451239 -0.22967941      NA 0.27873598
## [21] -1.52801554      NA      NA      NA -1.52026283
## [26]      NA 1.61093385 0.36989951      NA 0.04352648
## [31]      NA      NA -0.05411356 -1.53893110      NA
## [36]      NA 0.42847048 0.84603106
```

A shorthand way of specifying multiple negative numbers is to put the negative sign out in front of the vector of positive numbers. Type `x[-c(2, 10)]` to get the exact same result.

```
x[-c(2, 10)]
```

```
## [1]      NA 1.0422525      NA      NA -0.09266823
## [6] -0.50033332 -1.40475937      NA -0.33069581 -1.24412832
## [11]      NA      NA      NA -0.40215828      NA
## [16]      NA -1.49451239 -0.22967941      NA 0.27873598
## [21] -1.52801554      NA      NA      NA -1.52026283
## [26]      NA 1.61093385 0.36989951      NA 0.04352648
## [31]      NA      NA -0.05411356 -1.53893110      NA
## [36]      NA 0.42847048 0.84603106
```

So far, we've covered three types of index vectors – logical, positive integer, and negative integer. The only remaining type requires us to introduce the concept of 'named' elements. Create a numeric vector with three named elements using `vect <- c(foo = 11, bar = 2, norf = NA)`.

```
vect <- c(foo = 11, bar = 2, norf = NA)
```

We can also get the names of `vect` by passing `vect` as an argument to the `names()` function. Give that a try.

```
names(vect)
```

```
## [1] "foo" "bar" "norf"
```

Alternatively, we can create an unnamed vector `vect2` with `c(11, 2, NA)`. Do that now.

```
vect2 <- c(11, 2, NA)
```

Then, we can add the `names` attribute to `vect2` after the fact with `names(vect2) <- c("foo", "bar", "norf")`. Go ahead.

```
names(vect2) <- c("foo", "bar", "norf")
```

Now, let's check that `vect` and `vect2` are the same by passing them as arguments to the `identical()` function.

```
identical(vect, vect2)
```

```
## [1] TRUE
```

Now, back to the matter of subsetting a vector by named elements. Which of the following commands do you think would give us the second element of `vect`?

```
print("3: vect['bar']")
```

```
## [1] "3: vect['bar']"
```

Likewise, we can specify a vector of names with `vect[c("foo", "bar")]`. Try it out.

```
vect[c("foo", "bar")]
```

```
## foo bar
```

```
## 11  2
```